Listing 1: Annotated configuration for Graph-Augmented Neural Cellular Automata (Distill-style premultiplied RGBA loss)

```json
{
  "data": {
    "emojis_dir": "data/emojis",                // Folder containing RGBA PNG assets used as targets
    "targets": [ "bacteria.png","heart.png","gecko.png","tooth.png",
                 "star.png","smile.png","four_leaf_clover.png","squid.png",
                 "eye.png","real_gecko.png" ],  // All available target filenames
    "active_target": "gecko.png",               // Target actually used for training/testing
    "img_size": 40                              // Canvas size (HxW) in cells/pixels
  },

  "model": {
    "n_channels": 16,                           // Total channels = RGBA(4) + hidden(12)
    "perception": { "sobel": true, "identity": true }, // Fixed depthwise filters: identity + Sobel
    "update_mlp": { "hidden_dim": 128, "layers": 2 },  // 1x1 conv MLP width and depth
    "layer_norm": false,                        // (Unused here; GroupNorm is applied on dx)
    "fire_rate": 0.5,                           // Default firing prob at test (training uses range below)
    "update_gain": 0.06,                        // Step size after tanh(dx) bounding
    "alpha_thr": 0.12,                          // Alive threshold on alpha (after 3x3 max-pool) for dynamics
    "use_groupnorm": true                       // Apply GroupNorm(1,C) to dx before bounding
  },

  "training": {
    "pool_size": 1024,                          // Size of SamplePool (diversity/memory of states)
    "batch_size": 16,                           // Mini-batches per epoch
    "steps_per_epoch": 800,                     // Steps per epoch
    "nca_steps_min": 48,                        // Short rollout length range (min)
    "nca_steps_max": 80,                        // Short rollout length range (max)
    "long_rollout_prob": 0.4,                   // Chance to use a long rollout for a batch
    "long_rollout_steps_min": 200,              // Long rollout length (min)
    "long_rollout_steps_max": 400,              // Long rollout length (max)
    "fire_rate_min": 0.5,                       // Training: lower bound for random firing prob
    "fire_rate_max": 0.9,                       // Training: upper bound for random firing prob

    "num_epochs": 1000,                         // Planned number of epochs
    "learning_rate": 0.0002,                    // Adam step size
    "optimizer": "Adam",                        // Optimizer type
    "weight_decay": 1e-5,                       // L2 regularization
    "gradient_clip": 1.0,                       // Max global grad-norm (stability)

    "loss": "premult_rgba_mse",                 // Full-canvas MSE on premultiplied RGBA
    "loss_alpha_thr": 0.12,                     // (Unused by premult RGBA loss) kept for masked-loss ablations
    "loss_lam_area": 0.0,                       // (Unused with premult RGBA) keep 0
    "stability_enabled": true,                  // Optional: K-step stability rollouts for near-target states
    "stability_K": 48,                          // Stability rollout length (used only if enabled in code)
    "stability_threshold": 0.03,                // If per-sample loss < this, include in stability phase
    "stability_weight": 0.65,                   // Weight of stability loss term
```

```
    "reset_worst_prob": 0.10,                     // Fraction of worst samples in batch to reseed each step
    "random_reseed_prob": 0.05,                   // Chance to reseed one random sample (prevents collapse)
    "loss_lam_bg_alpha": 0.0,                     // (Unused with premult RGBA) keep 0
    "loss_lam_bg_rgb": 0.0,                       // (Unused with premult RGBA) keep 0

    "scheduler": {                               // LR schedule
      "type": "StepLR",                          // StepLR: decay LR by gamma every step_size epochs
      "step_size": 150,
      "gamma": 0.85
    }
  },

  "logging": {
    "checkpoint_interval_epochs": 5,             // Save model every N epochs
    "log_interval": 500,                         // Print/train-scalar interval (steps)
    "save_interval": 1000,                       // Image save interval (steps)
    "results_dir": "outputs",                    // Root folder for images/diagnostics
    "checkpoint_dir": "outputs/checkpoints",     // Unused by current trainer path mapping
    "visualize_interval": 1000                   // Add image to TensorBoard every N steps
  },

  "misc": {
    "device": "cuda",                            // "cuda" or "cpu"
    "seed": 42                                   // RNG seed for reproducibility
  },

  "graph_augmentation": {
    "d_model": 16,                               // Dim of Q/K projections used for attention over offsets
    "attention_radius": 4,                       // Max |dx|,|dy| for sampled mid-range neighbor offsets
    "num_neighbors": 8,                          // Offsets sampled each step (from the radius ring)
    "gating_hidden": 32,                         // Width of optional channel-wise gate MLP
    "message_gain": 0.25,                        // Scale after tanh on graph message (residual strength)
    "hidden_only": true,                         // If true, graph message updates hidden channels only
    "message_rate": 0.2,                         // Probabilistic use of graph message when msg_every=1
    "message_every": 3                           // Temporal sparsity: use graph every N steps (if >1)
  },

  "damage": {
    "start_epoch": 100,                          // Start applying damage as a curriculum from this epoch
    "prob": 0.3,                                 // Global chance that a batch sees damage this step
    "per_sample_prob": 0.4,                      // Within a damaged batch, fraction of samples damaged
    "kinds": {                                   // Mixture weights of damage types
      "square": 0.35, "circle": 0.25, "stripes": 0.10,
      "alpha_drop": 0.15, "saltpepper": 0.05, "gaussian": 0.10
    },
    "size_min": 6,                               // Damage size range (pixels/cells)
    "size_max": 18,
    "stripe_width": 6,                           // Stripe thickness for striped damage
    "alpha_thr": 0.2,                            // Alive threshold used by damage ops that depend on alpha
```

```json
    "alpha_dropout_p": 0.15,                   // Probability for channelwise alpha dropout damage
    "salt_pepper_p": 0.02,                     // Probability per pixel for salt-and-pepper noise
    "gaussian_softness": 0.35,                 // Edge softness for Gaussian/blurred damages
    "hidden_noise_sigma": 0.0                  // Std dev of noise injected into hidden channels (if any)
  }
}
```